

array vs NumPy, Creation, Broadcasting, Masking & Stats — Q&A; (Set 22)

Q1. Benefits of the built-in array package

array.array provides compact, memory-efficient, C-contiguous storage for a single primitive type (ints, floats, etc.), supports fast binary I/O (frombytes/tobytes), and avoids per-element Python object overhead.

Q2. Limitations of the array package

Single fixed typecode; mostly 1D; few numerical operations (no ufuncs/broadcasting); no views/strides; much smaller API than NumPy.

Q3. Main differences: array vs numpy

array: stdlib, 1D homogeneous buffers, minimal math.

NumPy: external package, N-D arrays, rich dtypes (float, int, bool, datetime, complex, fixed-len strings, object), vectorized ufuncs, broadcasting, views/strides, BLAS/LAPACK integration.

Q4. Distinctions between empty, ones, zeros (NumPy)

np.empty(shape): allocates uninitialized memory (contents arbitrary, fastest).

np.ones(shape): allocates and fills with 1s.

np.zeros(shape): allocates and fills with 0s. Use empty when you will overwrite every element.

Q5. Role of the callable in np.fromfunction

The callable `f` is invoked with index grids (arrays of coordinates for each axis); its return value populates the array. Example: `np.fromfunction(lambda i,j: i+j, (3,3), dtype=int)`.

Q6. Array + scalar ($A + n$)

NumPy broadcasts the scalar across all elements and returns an elementwise sum array; dtype promotion rules apply.

Q7. Combined op=assign with scalars ($+=$, $*=$)

Yes. `A += n` or `A *= n` performs in-place broadcasting where possible (mutates `A`). If `n` cannot be cast safely to `A.dtype` with the current casting rule, NumPy may raise an error.

Q8. Fixed-length strings in NumPy

Yes (dtype='Sk' for bytes or 'Uk' for Unicode, fixed length `k`). Assigning a longer string is truncated to fit. Object dtype ('O') stores arbitrary Python strings without truncation.

Q9. Combining two NumPy arrays with + or *; conditions

Operations are elementwise with broadcasting. Shapes must be equal or broadcast-compatible: for each trailing dimension, sizes must match or one of them must be 1.

Q10. Best way to use a Boolean array to mask another array

Use boolean indexing: $B = A[\text{mask}]$ to select elements; $A[\text{mask}] = \text{value}$ to assign. mask must be same shape as A (or match along the indexed axis).

Q11. Three ways to get standard deviation; fastest → slowest

- 1) NumPy: `np.std(arr)` (C-level vectorized)
 - 2) statistics module: `statistics.pstdev(arr)` / `statistics.stdev(arr)` (pure Python but optimized with `math.fsum`)
 - 3) Manual Python loop or two-pass with `sum/math.fsum`.
- Pandas `Series.std()` typically wraps NumPy and is comparable to #1.

Q12. Dimensionality of a Boolean mask-generated array

Using $A[\text{mask}]$ where mask has the same shape returns a 1-D array of the selected elements. If you mask along an axis (e.g., $A[\text{mask}, :]$), that axis is reduced; remaining axes are preserved.